

Documentation

SVM (classification):

General info on the dataset used:

- Name of the dataset: seed_dataset
- Number of classes: 2
- Labels are: Çerçvelik and Ürgüp Sivrisi.
- Number of samples: 2500
- Number of samples in training is 1750.
- Number of samples in validation is 500.
- Number of samples in testing is 750.

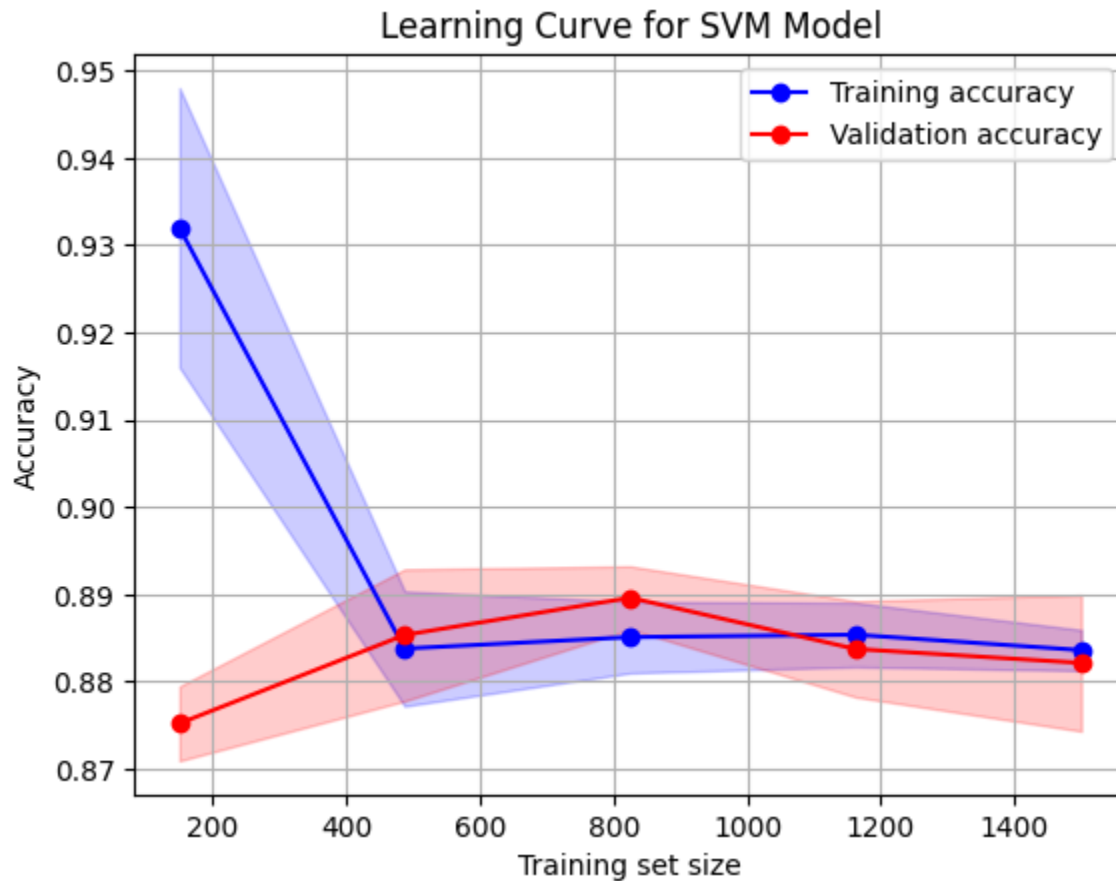
Implementation details:

- At feature extraction phase, 8 features are extracted, they are: Area, Perimeter, Major_Axis_Length, Minor_Axis_Length, Solidity, Extent, Roundness, Compactness.
- Cross-Validation **is** used in the model, and the number of fold is 5, and the ratio of training/validation is 4/1.
- Hyperparameters used in our model is:

- C Hyperparameter which controls the regularization strength
- Kernel which specifies the type of kernel used in the SVM model(linear).
- K which determines the number of subsets in which the data is divided for cross validation
- Shuffle which determines whether to shuffle the data before splitting it into folds. It's (true)→ data will be shuffled.
- Random state which is used to ensure that the same randomization is used each time you run the code.


Results details:

Learning curve: It shows how error changes as the training set size increases



Accuracy: 0.872


Confusion Matrix: Visualizes the performance of the algorithm



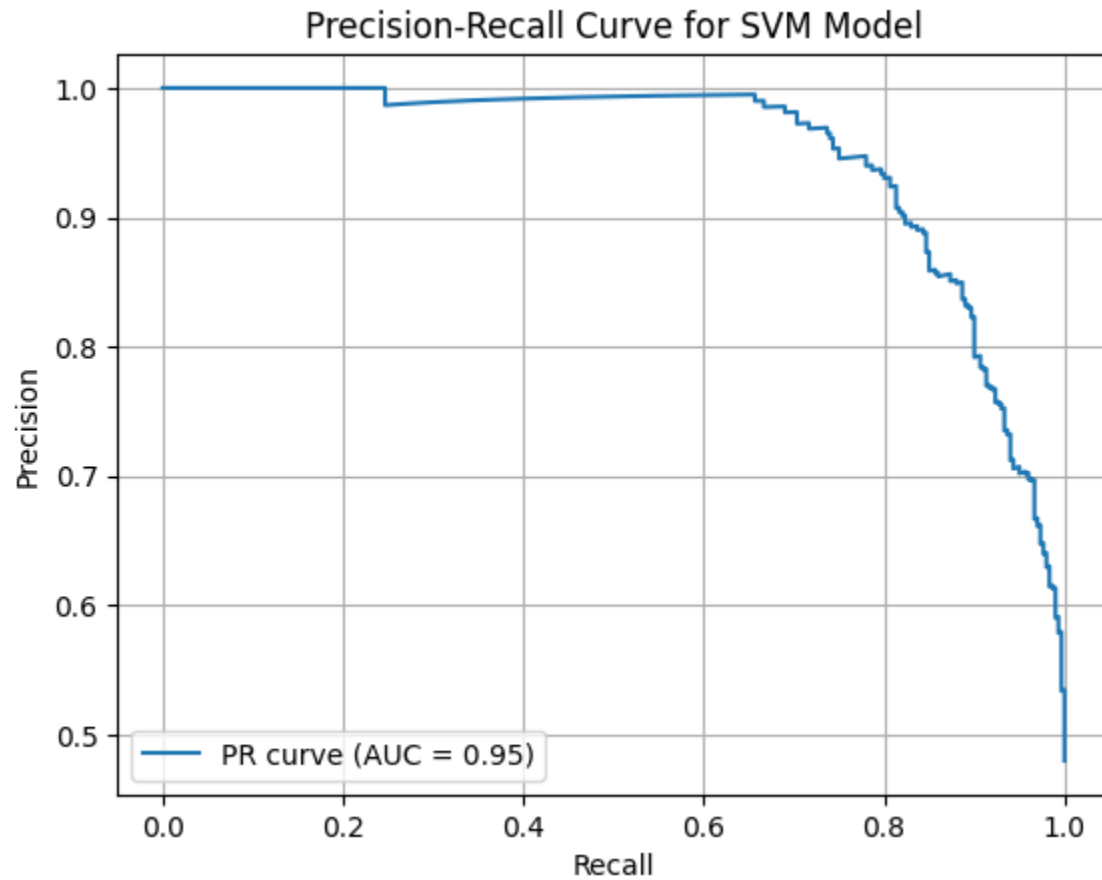
Classification Report: Represents the main metrics

	precision	recall	f1-score	support
0.0	0.86	0.91	0.88	325
1.0	0.89	0.83	0.86	300
accuracy			0.87	625
macro avg	0.87	0.87	0.87	625
weighted avg	0.87	0.87	0.87	625

Receiver Operating Characteristic (ROC) curve: It shows the connection between sensitivity and specificity, so it basically tells us how much the model is capable of distinguishing between classes



Precision-Recall curve: It shows the connection between precision and recall and, it provides insights into the model's ability to correctly classify the positive instances



Decision Tree (classification):

General info on the dataset used:

- Name of the dataset: seed_dataset
- Number of classes: 2
- Labels are: Çerçvelik and Ürgüp Sivrisi
- Number of samples: 2500
- Number of samples in training: 1750
- Number of samples in validation: 375
- Number of samples in testing: 375
- Number of features extracted: 9

Implementation details:

- Extracted features: ['Area', 'Perimeter', 'Major_Axis_Length', 'Minor_Axis_Length', 'Solidity', 'Extent', 'Roundness', 'Compactness', 'Class']
- Cross-Validation Scores: [0.81714286 0.82571429 0.84857143 0.82857143 0.82285714]
- Mean CV Score: 0.8285714285714285
- Number of Folds: 5
- Training Ratio (per fold): 0.8
- Validation Ratio (per fold): 0.2
- Hyperparameters used in our model is:

Criterion: gini

Max Depth: None

Min Samples Split: 2

Min Samples Leaf: 1

Splitter: best

Results details:

```
Decision Tree Model Accuracy (Test Set): 0.8426666666666667
```

```
Decision Tree Hyperparameters:  
Criterion: gini  
Max Depth: None  
Min Samples Split: 2  
Min Samples Leaf: 1  
Splitter: best  
Cross-Validation Scores: [0.81714286 0.82571429 0.84857143 0.82857143 0.82285714]  
Mean CV Score: 0.8285714285714285  
Number of Folds: 5  
Training Ratio (per fold): 0.8  
Validation Ratio (per fold): 0.2  
Decision Tree Model Accuracy (Test Set): 0.8426666666666667
```

```
Number of features extracted: 9
```

```
Extracted features: ['Area', 'Perimeter', 'Major_Axis_Length', 'Minor_Axis_Length', 'Solidity', 'Extent', 'Roundness', 'Compactness', 'Class']
```


Neural Network (Regression):

General info on the dataset used:

- Name of the dataset: vehicles
- Label: Price.
- Number of samples: 426,880
- Number of samples in training is 298,816.
- Number of samples in testing is 128,064.

Implementation details:

- At feature extraction phase: 12 features are extracted, they are: manufacturer, Model, condition, cylinders, fuel, odometer, title_status, transmission, drive, type, paint_color, and car_age.
- Model Architecture Hyperparameters:
 - Number of hidden layers: 2
 - Number of units in each hidden layer: 64
 - Activation function for hidden layers: ReLU (Rectified Linear Unit)
 - Output layer activation: Linear
- Training Hyperparameters:
 - Loss function: Mean Absolute Error (MAE)
 - Optimizer: RMSprop
 - Learning rate: 0.02
 - Number of epochs: 200
 - Batch size: 16

Results details:

```
test_mae_score
```



3535.3828125

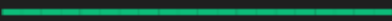
```
test_mse_score
```



3535.090087890625

```
y_pred = model.predict(X_test)
y_pred[0]
```



2410/2410  2s 685us/step

array([7882.9565], dtype=float32)

```
r2_accuracy = metrics.r2_score(y_test, y_pred)
print("Accuracy : ", r2_accuracy)
```



Accuracy : 0.8061071663159987